

NLP Recent History

Text Features

- Natural Language Processing = Machine Learning with Text
- Step 1 – Identify text-based features
- Step 2 – Label data for a task
- Step 3 – Train a classifier

Text Features – Old School

- Lexical Features
 - Part-of-speech tags (Verb/Noun/Adj)
 - Stems / Lemmas
 - Word Frequency
 - Psycholinguistic Features
 - Terminologies
 - Morphology (affixes)

Text Features – Old School

- Sentential Features
 - Sentence Structure
 - Parse Tree Features
 - Sentence Dependencies
 - POS ratios
 - Punctuation frequency
 - Averages of lexical features

Text Features – Old School

- Document Features
 - Bag-of-words
 - TF-IDF Vectors

Text Features – Word Embeddings

- Represent each word as a vector of numbers

Text Features – Word Embeddings

- Represent each word as a vector of numbers

The cat sat on the mat

Text Features – Word Embeddings

- Represent each word as a vector of numbers

the →

0.23
0.12
0.65
...
...

The cat sat on the mat

Text Features – Word Embeddings

- Represent each word as a vector of numbers



Text Features – Word Embeddings

- Represent each word as a vector of numbers

$\overrightarrow{\text{the}}$	$\overrightarrow{\text{cat}}$	$\overrightarrow{\text{sat}}$	$\overrightarrow{\text{on}}$	$\overrightarrow{\text{the}}$	$\overrightarrow{\text{mat}}$
0.23	0.38	0.31	0.43	0.72	0.44
0.12	0.12	0.02	0.82	0.91	0.10
0.65	0.79	0.63	0.61	0.23	0.31
...
...
The	cat	sat	on	the	mat

Text Features – Count-based Embeddings

- The cat sat on the mat
- Dogs and cats are animals
- Cats like milk
- The dog chased the stick
- Dogs like bones

	cat	sat	mat	dog	chase	stick	animal	bone	milk
cat									
sat									
mat									
dog									
chase									
stick									
animal									
bone									
milk									

Text Features – Count-based Embeddings

- The cat sat on the mat
- Dogs and cats are animals
- Cats like milk
- The dog chased the stick
- Dogs like bones

	cat	sat	mat	dog	chase	stick	animal	bone	milk
cat	3	1	1	1	0	0	1	0	1
sat	1	1	1	0	0	0	0	0	0
mat	1	1	1	0	0	0	0	0	0
dog	1	0	0	3	1	1	1	1	0
chase	0	0	0	1	1	1	0	0	0
stick	0	0	0	1	1	1	0	0	0
animal	1	0	0	1	0	0	1	0	0
bone	0	0	0	1	0	0	0	1	0
milk	1	0	0	1	0	0	0	0	1

Text Features – Count-based Embeddings

- The cat sat on the mat
- Dogs and cats are animals
- Cats like milk
- The dog chased the stick
- Dogs like bones

cat	sat	mat	dog	chase	stick	animal	bone	milk
3	1	1	1	0	0	1	0	1
1	1	1	0	0	0	0	0	0
1	1	1	0	0	0	0	0	0
1	0	0	3	1	1	1	1	0
0	0	0	1	1	1	0	0	0
0	0	0	1	1	1	0	0	0
1	0	0	1	0	0	1	0	0
0	0	0	1	0	0	0	1	0
1	0	0	1	0	0	0	0	1

Text Features – Count-based Embeddings

- The cat sat on the mat
- Dogs and cats are animals
- Cats like milk
- The dog chased the stick
- Dogs like bones

- Issue
 - Dimensionality = vocab size
- Reduce with PCA..

cat	sat	mat	dog	chase	stick	animal	bone	milk
3	1	1	1	0	0	1	0	1
1	1	1	0	0	0	0	0	0
1	1	1	0	0	0	0	0	0
1	0	0	3	1	1	1	1	0
0	0	0	1	1	1	0	0	0
0	0	0	1	1	1	0	0	0
1	0	0	1	0	0	1	0	0
0	0	0	1	0	0	0	1	0
1	0	0	1	0	0	0	0	1


Text Features – Count-based Embeddings

- Which animal likes milk?


animal		milk		cat
1		1		3
0		0		1
0		0		1
1	+	0	~	1
0		0		0
0		0		0
1		0		1
0		0		0
0		1		1

Text Features – Learnt Embeddings (Word2Vec)

- Learn a set of word vectors that can be used to predict surrounding words
- Initialise random vectors and train neural network, updating vector weights with network
- End result is a set of vectors that perform like count-based, but without counting




Text Features – Learnt Embeddings (Word2Vec)




$$\vec{\text{king}} - \vec{\text{man}} + \vec{\text{woman}} \approx \vec{\text{queen}}$$

Text Features – Learnt Embeddings (Word2Vec)



Text Features – Learnt Embeddings (Word2Vec)




Example Tasks

- Classification
 - Sentiment Analysis
 - Readability
 - Categorisation
- Text Generation
 - Machine Translation
 - Summarisation
 - Data-to-text
- Natural Language Understanding
 - Inference / Reasoning
 - Chatbots / Intelligent Tutoring

<http://nlpprogress.com/>

Sequence classification

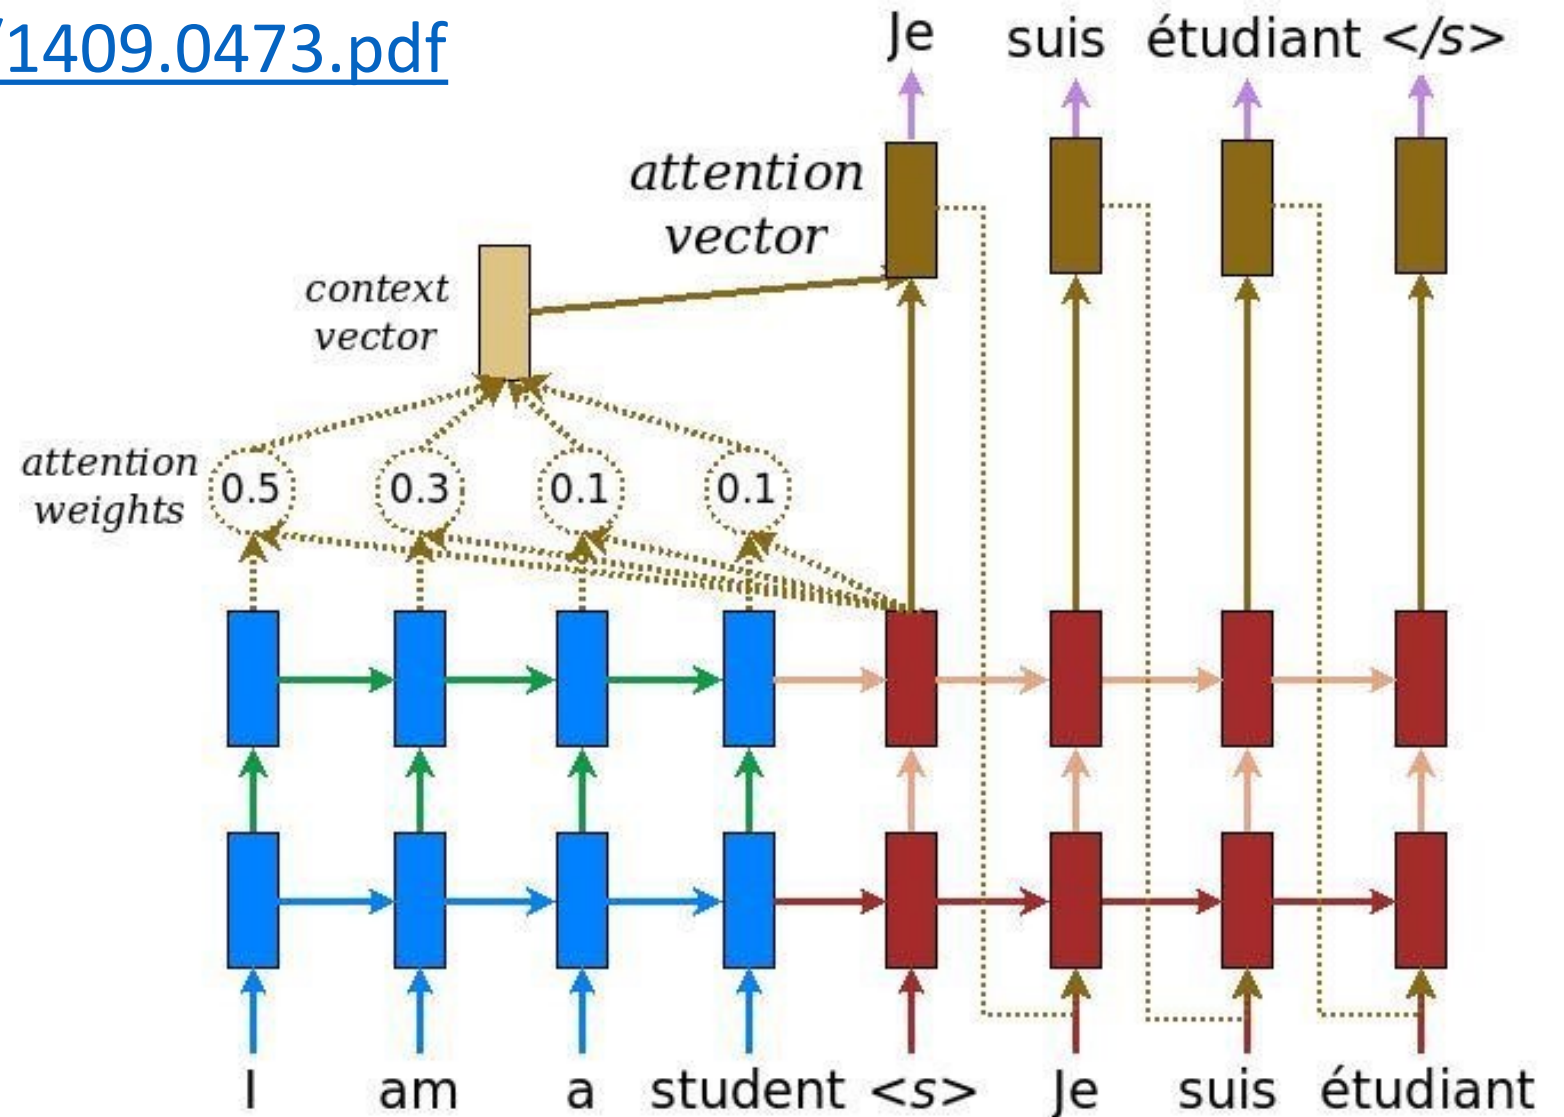
- Long Short Term Memory



Sequence classification

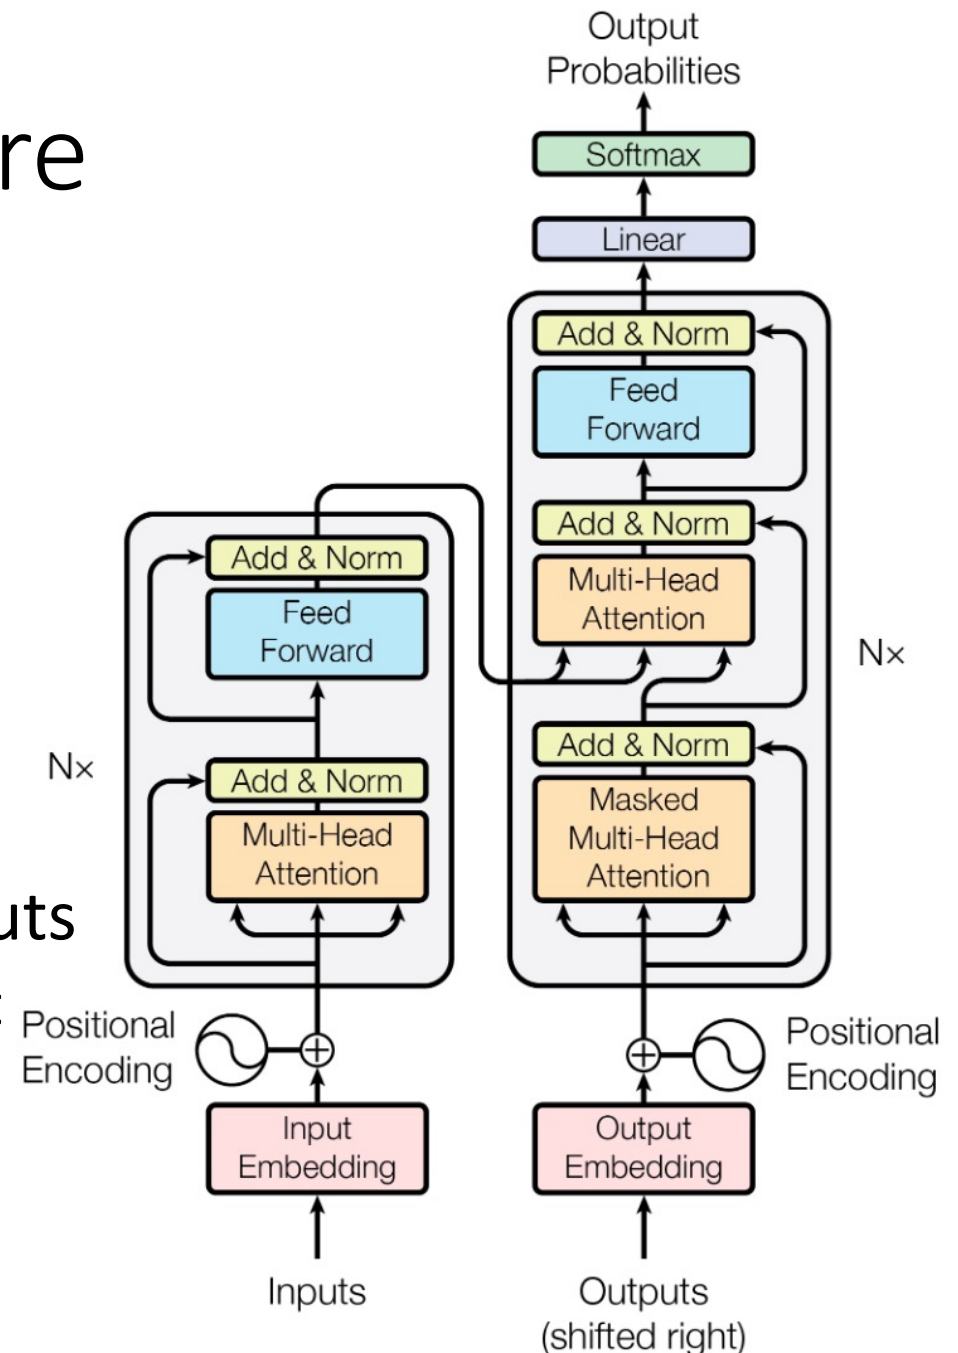
<https://arxiv.org/pdf/1409.0473.pdf>

- LSTMs + Attention



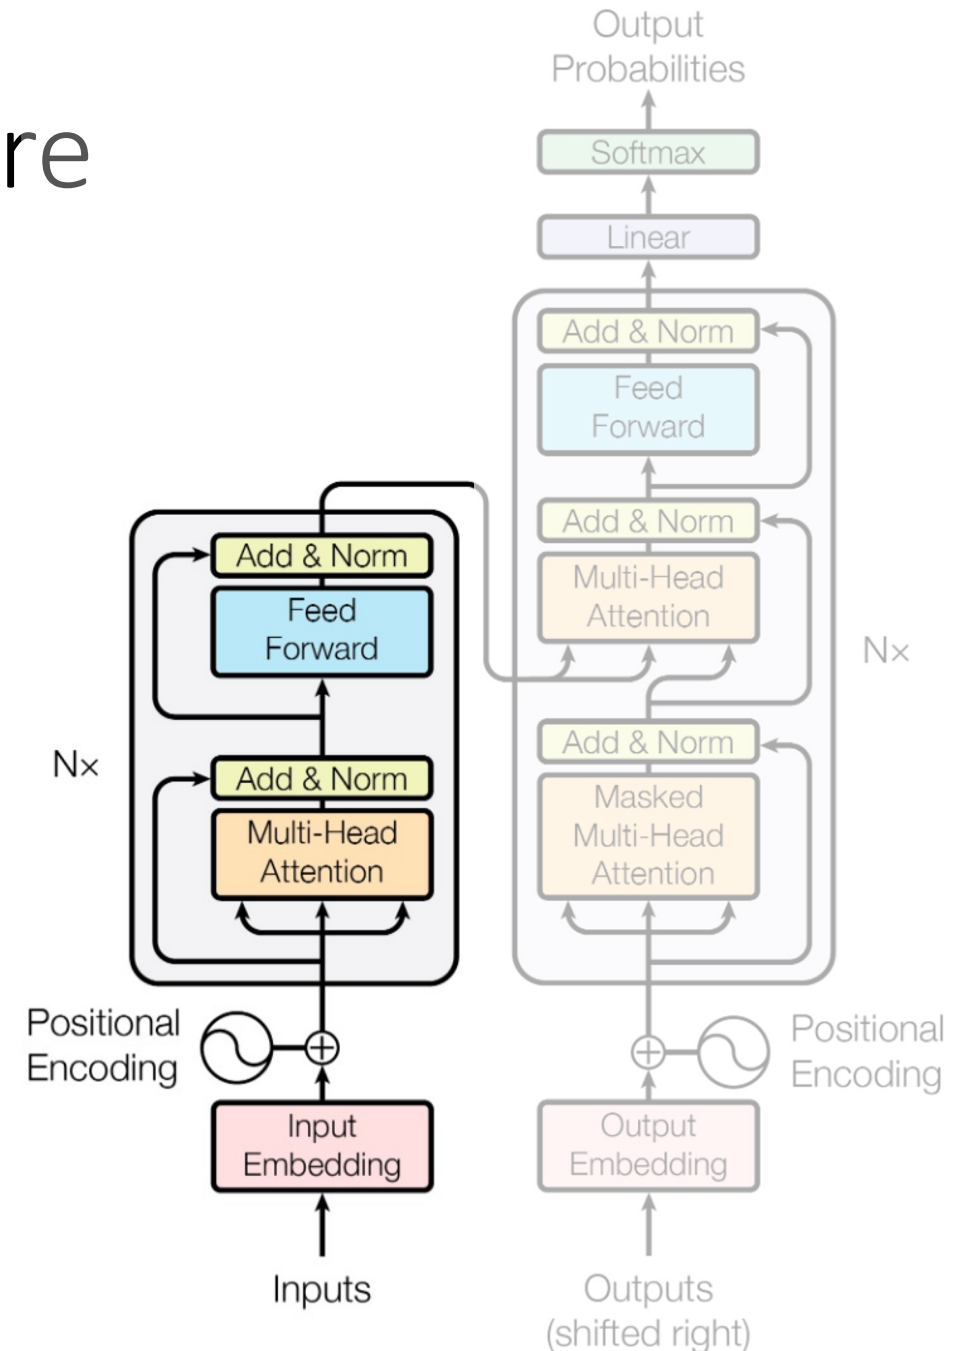
The Transformer - Architecture

- Encoder-Decoder architecture
 - Feature space acts as an interlingua
- Encoder
 - Inputs (words) are encoded using embeddings
 - Transforms input embeddings to feature space
 - 6 stacked layers attention and feed-forward
- Decoder – transform feature space to outputs
 - Converts output of last encoder layer into a set of probabilities which predict the next word



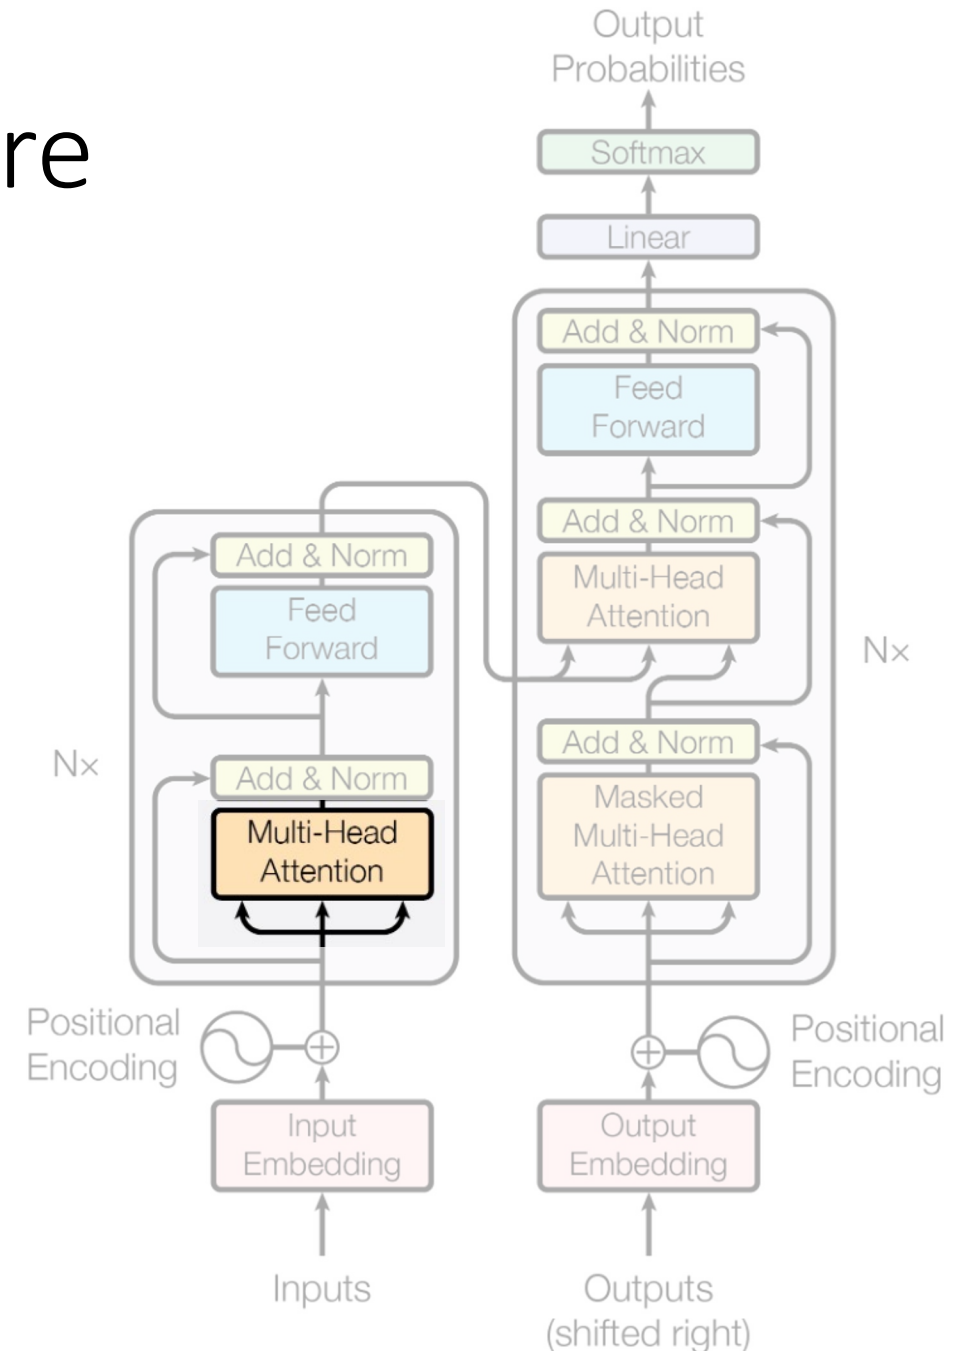
The Transformer - Architecture

- Encoder Layer
- Creates context vector for input sequence
- N=6 stacked layers



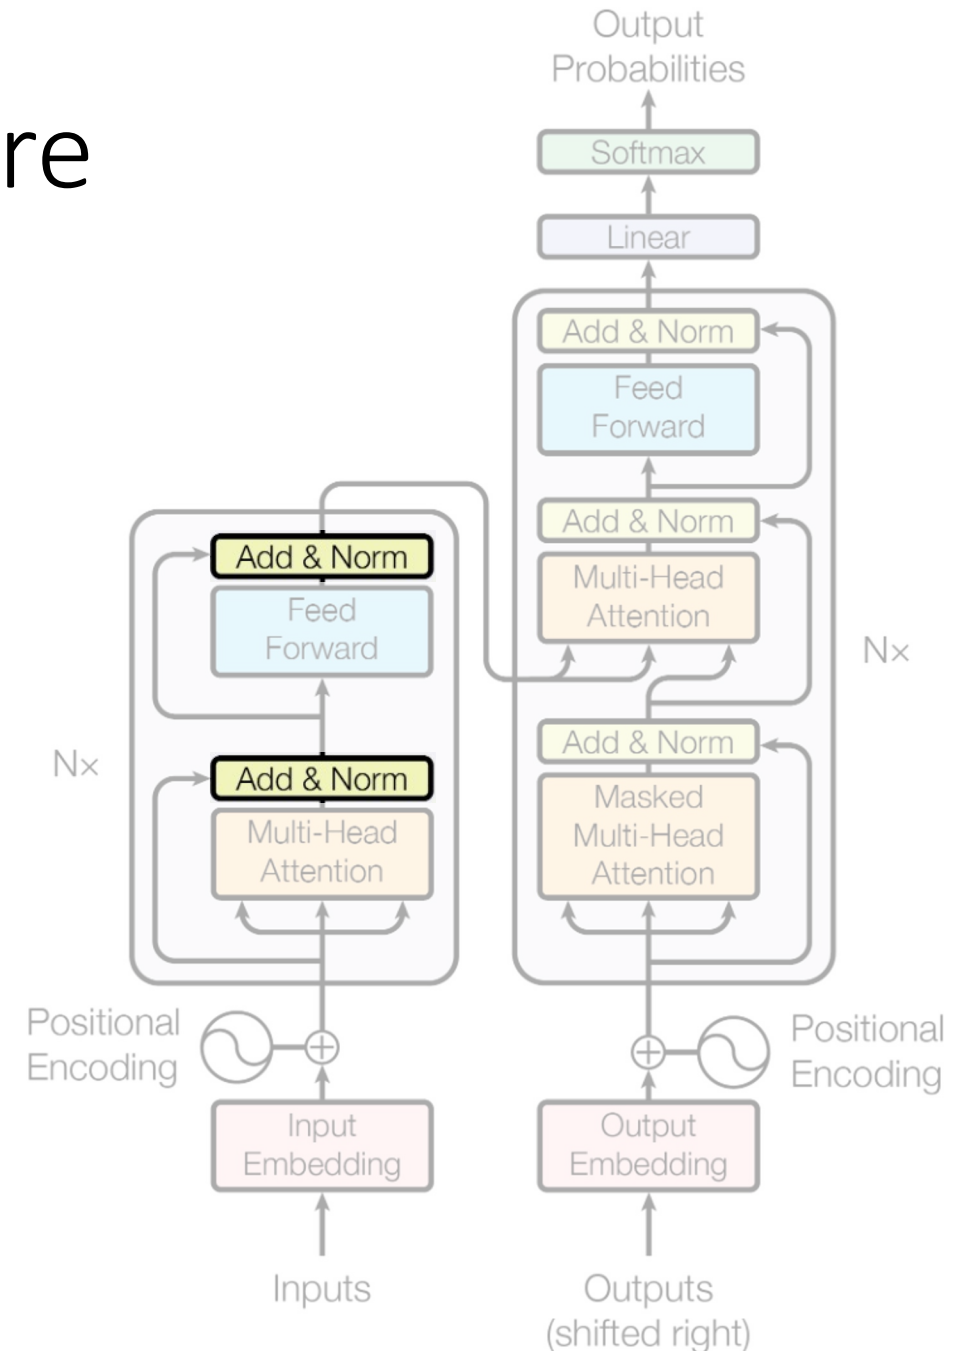
The Transformer - Architecture

- Encoder – multi-headed self attention
- Self-attention
 - All inputs available to encoder
 - Attention selects relevant parts of input
 - Key, Value, Query (read paper for details)
- Multi-headed
 - Ran in parallel $N=8$ times
 - Output of each 'head' is concatenated



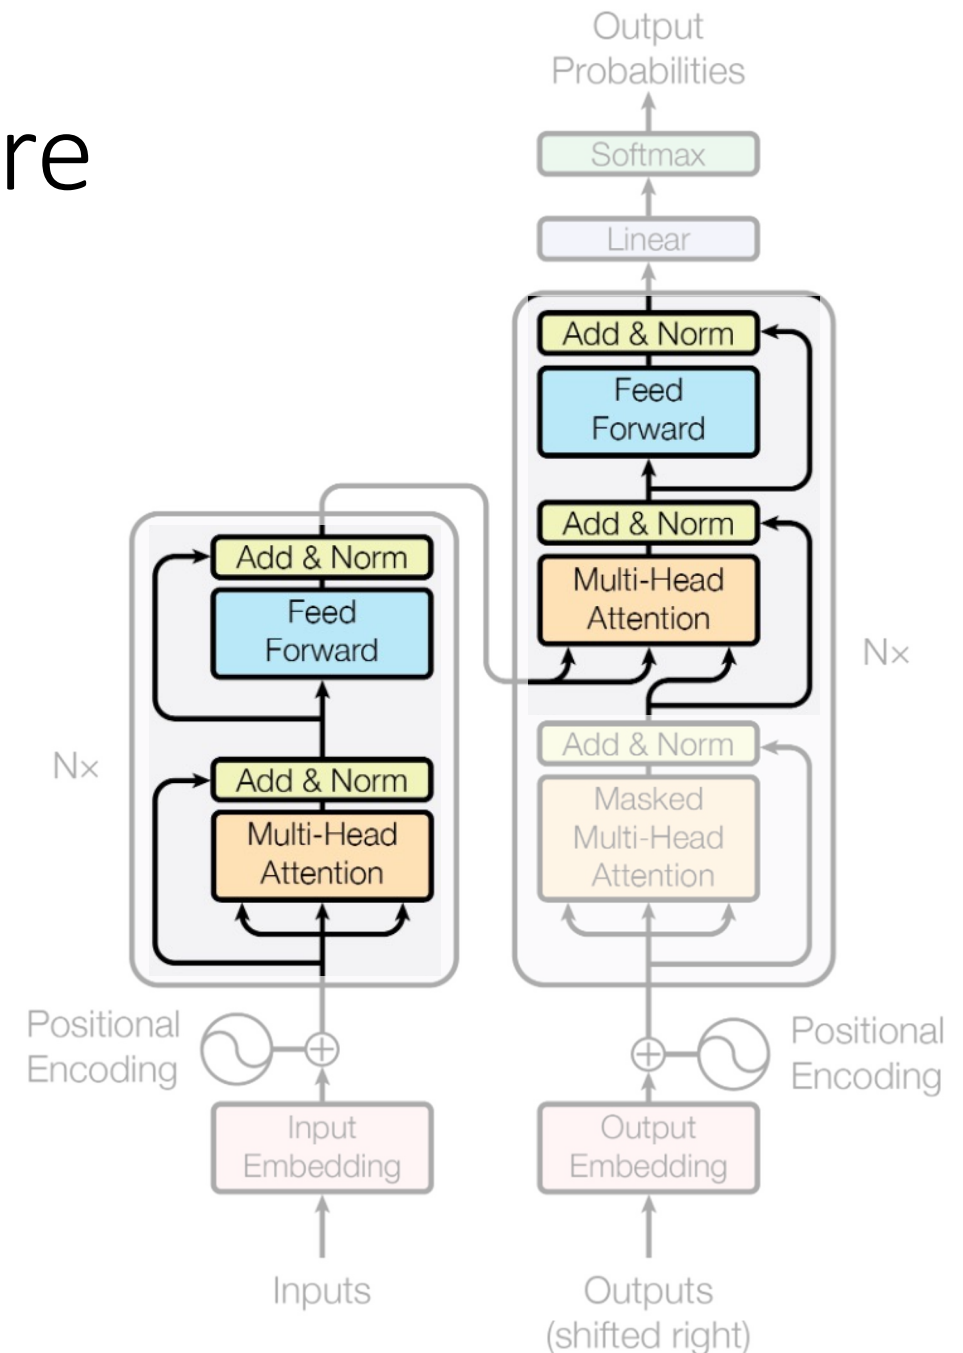
The Transformer - Architecture

- Encoder - Add and normalize
- Note that the input from before each step is carried through to the next step
- Normalisation prevents exploding gradient



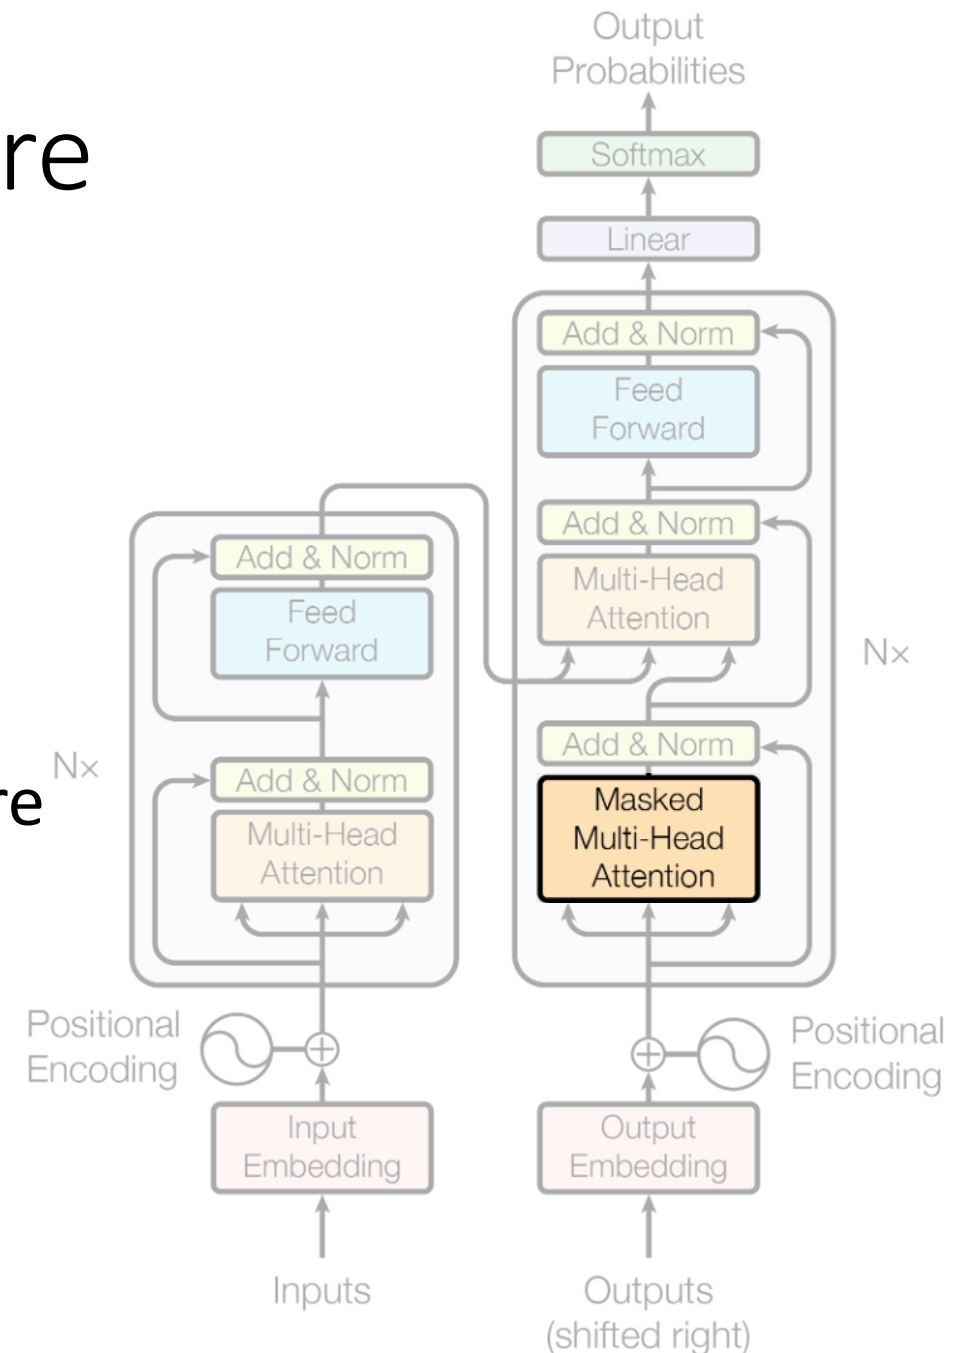
The Transformer - Architecture

- Decoder – Replicates encoder
- Multi-head self attention and feed forward network are carried over
 - Applied to the attention weights from encoder
 - Also applied to the last output from decoder



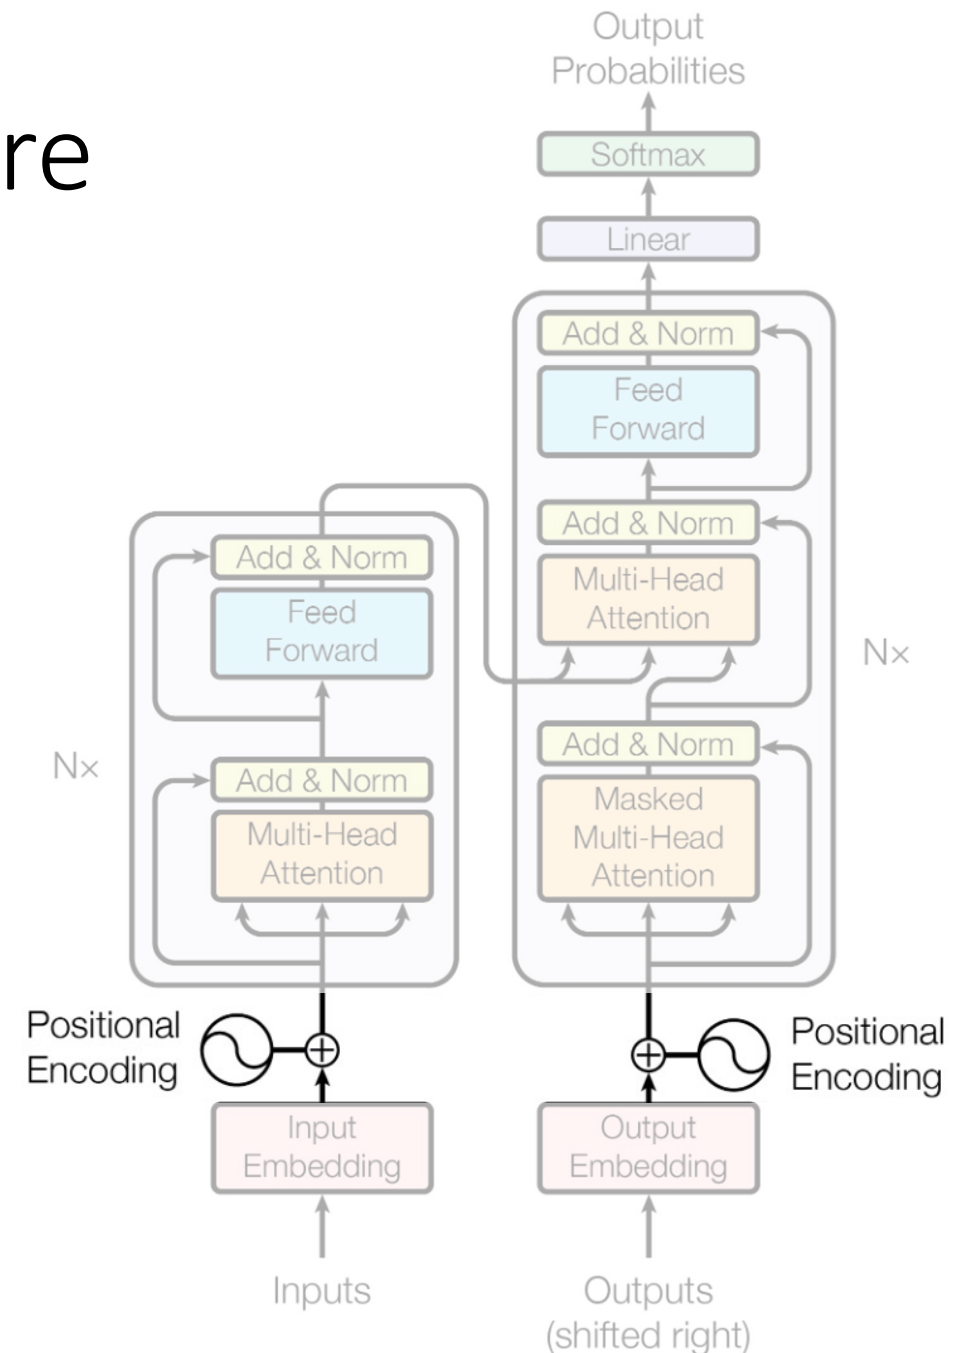
The Transformer - Architecture

- Decoder – Masked Multi-head attention
- Masking prevents information from future becoming available at decode time
- Otherwise, self-attention the same as before



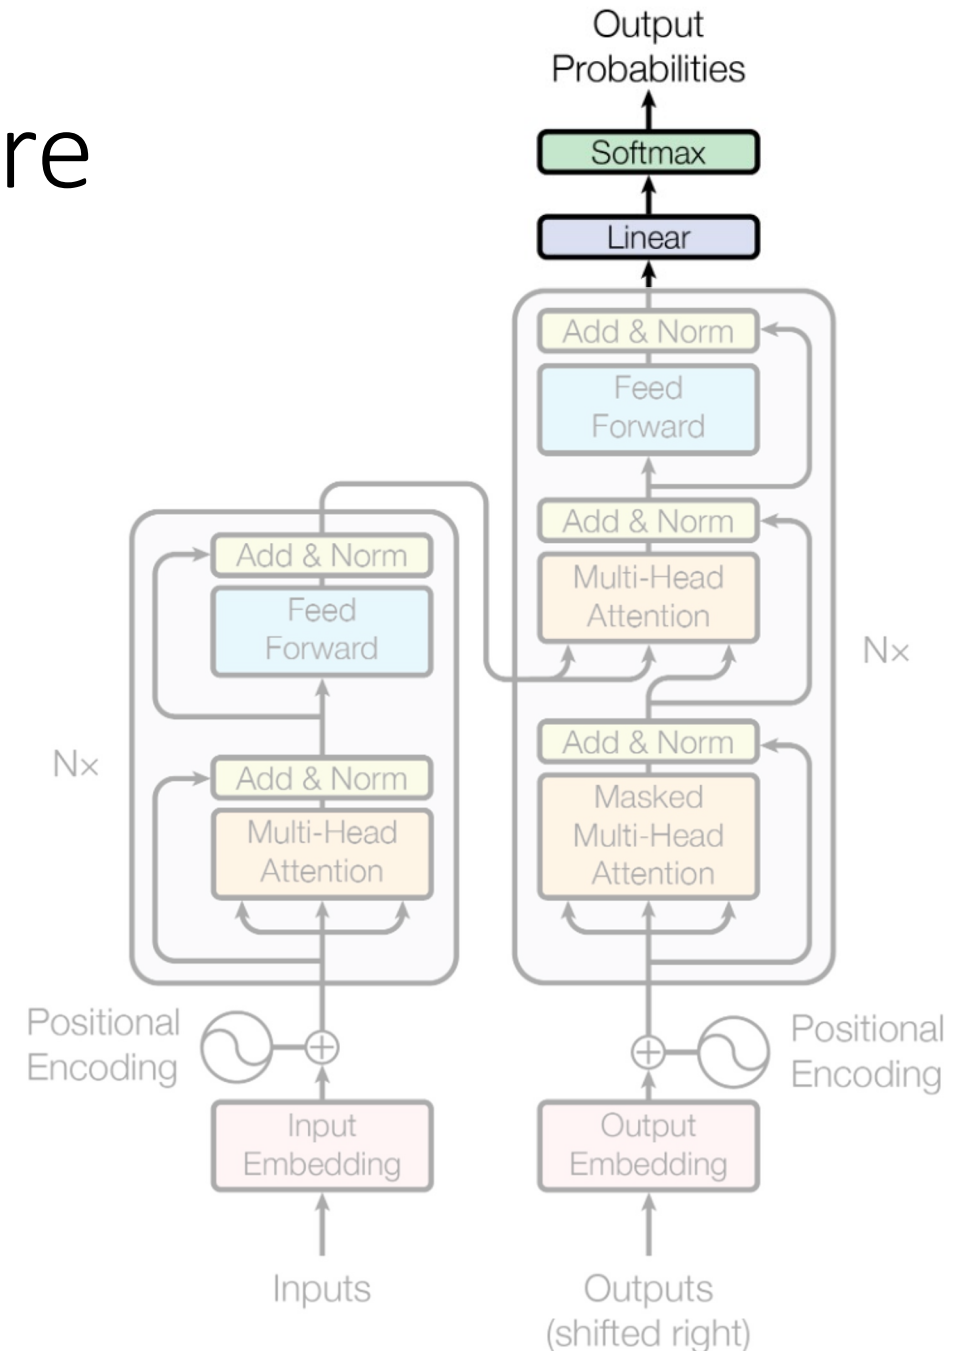
The Transformer - Architecture

- Encoder-Decoder – Sinusoidal Positional Encoding
- Encodes whereabout in the sentence the word is likely to have occurred
 - Applies sin/cos function to alternate dimensions of the embedding
 - This is summed with the word embeddings
 - So embedding contains info on meaning + position in sentence

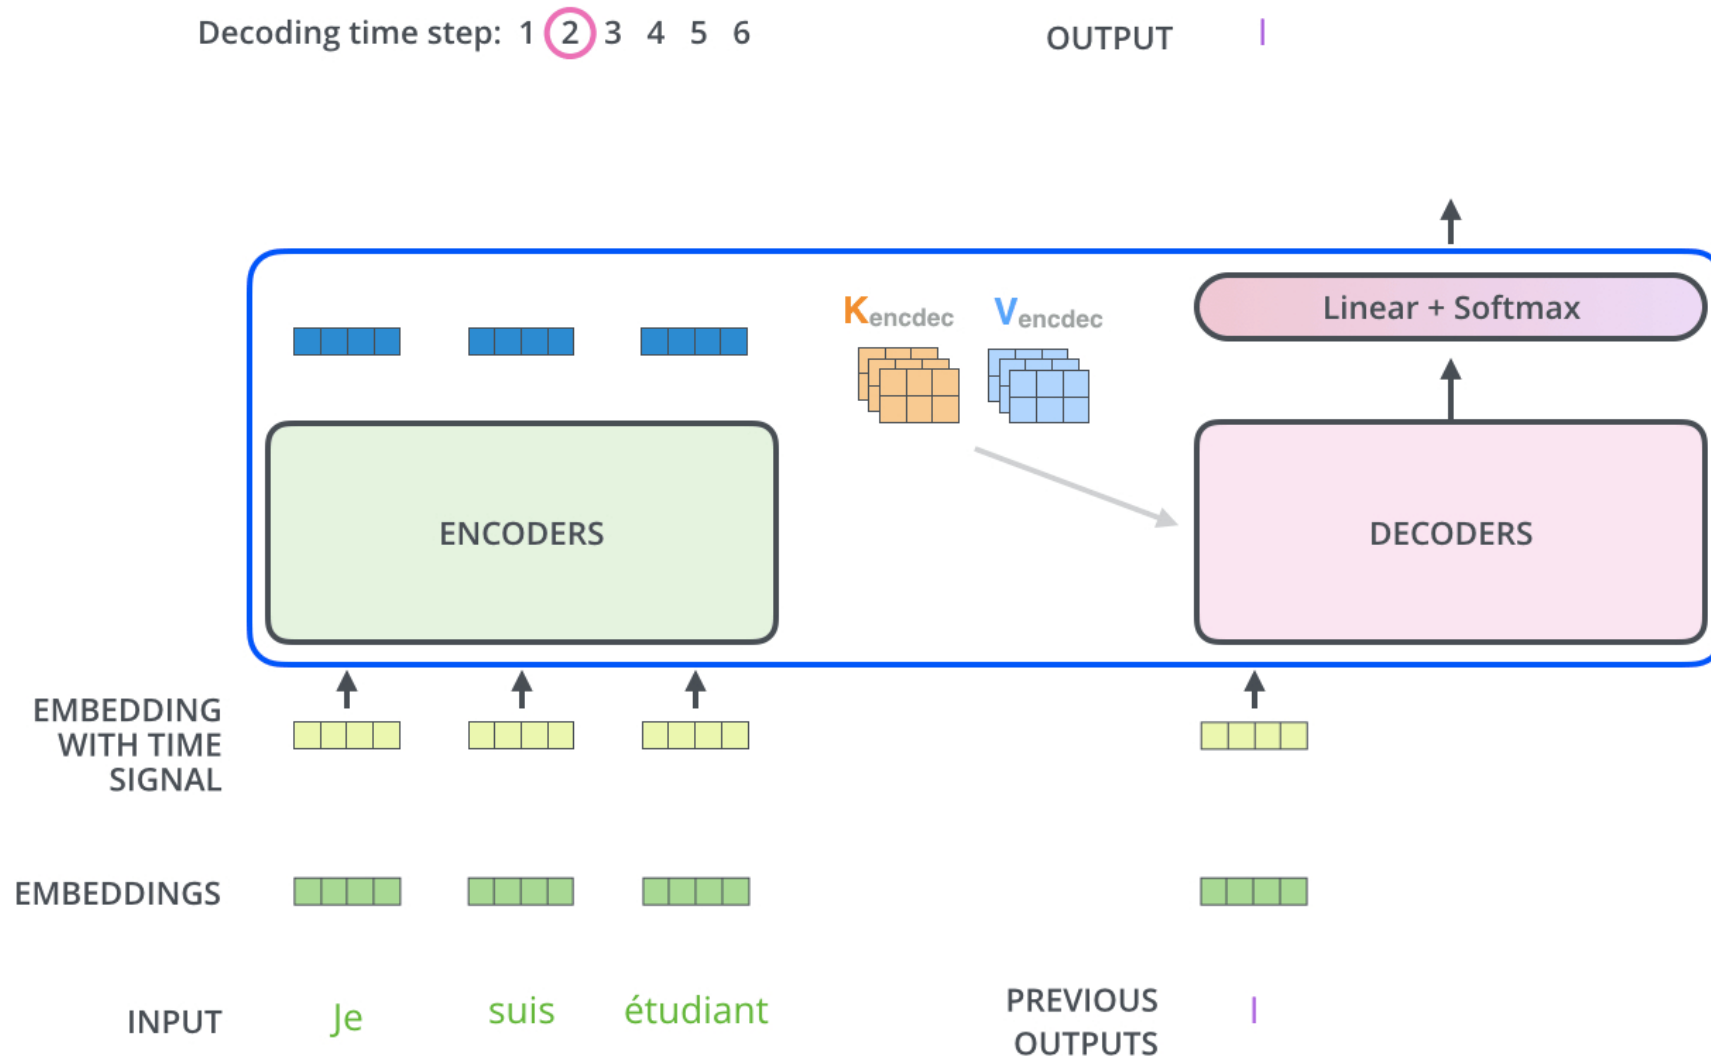


The Transformer - Architecture

- Output
- Linear layer
 - Fully connected NN
 - Converts k-dim vector to size of output vocab
- Softmax layer
 - Normalises probabilities in 0-1 range.




The Transformer - Visualisation



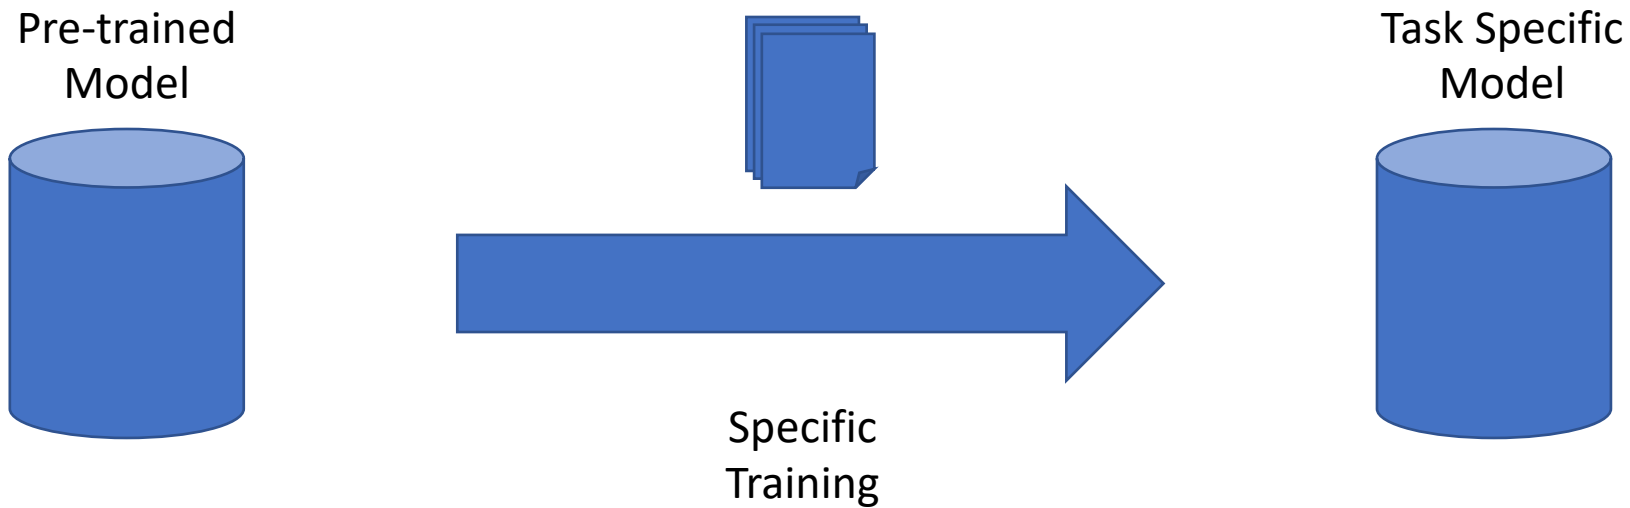
Pre-trained Models

- NLP would be very expensive if we had to retrain models every time
- Fortunately, pre-trained models are available that we can use



Fine-tuning

- Pre-trained models can be adapted to a specific task





Models

Datasets

Spaces

Docs

Solutions

Pricing







Log In


Sign Up


Tasks Libraries Datasets Languages Licenses
Other


Computer Vision


 Depth Estimation  Image Classification

 Object Detection  Image Segmentation



 Image-to-Image


 Unconditional Image Generation


 Video Classification



 Zero-Shot Image Classification



Natural Language Processing



 Text Classification  Token Classification



 Table Question Answering

 Question Answering

 Zero-Shot Classification  Translation

 Summarization  Conversational

 Text Generation  Text2Text Generation

 Fill-Mask  Sentence Similarity

Models 124,065

↑↓ Sort: Most Downloads

bert-base-uncased

 • Updated Nov 16, 2022 • ↓ 25.8M • ♥ 464
Fri, 16 Dec 2022 15:44:21 GMT

gpt2

 • Updated Dec 16, 2022 • ↓ 14M • ♥ 469

openai/clip-vit-large-patch14


 • Updated Oct 4, 2022 • ↓ 10.2M • ♥ 160

xlm-roberta-large

 • Updated Jun 27, 2022 • ↓ 10M • ♥ 64

distilbert-base-uncased

 • Updated Nov 16, 2022 • ↓ 9.88M • ♥ 127

 prajjwal1/bert-tiny

Updated Oct 27, 2021 • ↓ 9.65M • ♥ 25

distilbert-base-uncased-finetuned-sst-2-english

 • Updated Dec 5, 2022 • ↓ 7.06M • ♥ 137

**Hugging Face**

Search models, datasets

Sign Up

Tasks

Libraries

Datasets


Languages

Licenses


Other

Filter Tasks by name

Computer Vision

 Depth Estimation Image Classification Object Detection Image Segmentation Image-to-Image Unconditional Image Generation Video Classification Zero-Shot Image Classification

Natural Language Processing

 Text Classification Token Classification Table Question Answering Question Answering Zero-Shot Classification Translation Summarization Conversational Text Generation Text2Text Generation Fill-Mask Sentence Similarity

Natural Language Processing



Text Classification



Token Classification



Table Question Answering



Question Answering



Zero-Shot Classification



Translation



Summarization



Conversational



Text Generation



Text2Text Generation



Fill-Mask



Sentence Similarity

distilbert-base-uncased-finetuned-sst-2-english

 • Updated Dec 5, 2022 • ↓ 7.06M • ♥ 137